

Q1-Final

March 1, 2020

```
In [1]: import numpy as np
import cv2
import matplotlib.pyplot as plt
import random
from glob import glob
import time

# get_ipython().magic('matplotlib inline')

In [2]: def construct_H_matrix(x, xs):
'''
    Construct the correspondance matrix for estimating Homography

    Keyword Arguments:
        x -- Image points from the First image
        xs -- Image points from the Second Image

    Return Values:
        A -- Constructed correspondance matrix
'''
A = np.zeros((0, 9))
for i in range(len(x)):
    a = np.array([
        x[i][0], x[i][1], 1, 0, 0, 0, -xs[i][0]*x[i][0], -xs[i][0]*x[i][1], -xs[i][0],
        0, 0, 0, x[i][0], x[i][1], 1, -xs[i][1]*x[i][0], -xs[i][1]*x[i][1], -xs[i][1],
    ])
    A = np.concatenate((A, a))
return A

In [3]: def transform_points(H, x):
'''
    Given homography transforms points into the homography frame of reference

    Keyword Arguments:
        H -- 3*3 homography matrix
        x -- N*2 Points to be transformed

    Return Values:
```

```

        xs -- N*3 Transformed Points
    """
    # Convert points into homogeneous coordinates
    x_homogeneous = np.concatenate((x, np.ones((len(x), 1))), axis=1)

    # Calculate the transformed points and normalize them
    xs = H @ x_homogeneous.T
    xs[0, :] = xs[0, :] / xs[2, :]
    xs[1, :] = xs[1, :] / xs[2, :]
    xs[2, :] = xs[2, :] / xs[2, :]
    xs = xs.T

    return xs

```

```

In [4]: def calculate_transformation_error(H, x, xs):
    """
        Calculate the transformation error between the transformed points and the ground truth

        Keyword Arguments:
            H                -- 3*3 homography matrix
            x                -- N*2 points to be transformed
            xs               -- N*2 ground truth points

        Return Values:
            transformation_error -- Return the L2 norm of the error
    """
    # Convert points into homogeneous coordinates
    xs_homogeneous = np.concatenate((xs, np.ones((len(xs), 1))), axis=1)

    # Calculate the projected points
    transformed_coords = transform_points(H, x)

    # Calculate the projection error
    transformation_error = np.linalg.norm(xs_homogeneous - transformed_coords, axis=1)

    return transformation_error

```

0.1 RANSAC

- For better estimation of the homography matrix, we perform RANSAC and try to get the homography matrix with maximum inliers.
- We calculate the inliers based on the transformation error, which is required to be below a certain preset threshold that we choose.

```

In [5]: def RANSAC(img_coords_1, img_coords_2, num_points, max_iterations=30000, thresh=0.1):
    """
        Get the best estimate of the Homography Matrix using RANSAC

```

```

Keyword Arguments:
    img_coords_1  -- Coordinates of the first image after matching
    img_coords_2  -- Coordinates of the first image after matching
    num_points    -- Number of matched points
    max_iterations -- Maximum Number of iterations to run RANSAC for (default=20)
    thresh        -- Threshold to compute inliers (default=0.1)

Return Values:
    H              -- Estimated Homography matrix using RANSAC
'''
min_transform_error = 999999999
max_inliers = -999999999

# Best estimate of Projection matrix by far
_H = np.zeros((3, 3))

inliers = []

for i in range(max_iterations):

    # Randomly select 4 world points and the corresponding image points
    idx = random.sample(range(0, num_points), 4)
    x  = img_coords_1[idx]
    xs = img_coords_2[idx]

    # Perform DLT and get the Transformation Matrix
    H = DLT(x, xs)

    # Calculate projection error
    transformation_error = calculate_transformation_error(H, img_coords_1, img_coord

    inliers = np.sum(transformation_error<thresh)

    if inliers > max_inliers:
        max_inliers = inliers
        _H = H

    # Repeat for a maximum number of iterations

return _H

```

0.2 Estimating the Homography Matrix

- To estimate the homography matrix we perform a DLT like estimation
- $x_2 = H_{21} \cdot x_1$

-

$$\begin{bmatrix} wx_{21} \\ wx_{22} \\ w \end{bmatrix} = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ h_7 & h_8 & h_9 \end{bmatrix} \cdot \begin{bmatrix} x_{11} \\ x_{12} \\ 1 \end{bmatrix} \quad (1)$$

- We take the SVD of the correspondance matrix which we define in the following way
- $\begin{bmatrix} x_1 & y_1 & 1 & 0 & 0 & 0 & -x_2x_1 & -x_2y_1 & -x_2 \\ 0 & 0 & 0 & x_1 & y_1 & 1 & -y_2x_1 & -y_2y_1 & -y_2 \end{bmatrix}$
- We require at least 4 points in order to estimate the homography matrix as it has 8 degrees of freedom. The ninth degree is removed due to scale ambiguity.
- We minimize the error, by taking the eigen vector corresponding to the smallest eigen value and reshape that to be our homography matrix

```
In [6]: def DLT(x, xs):
        """
            Estimate Homography matrix using DLT

            Keyword Arguments:
                x -- Image points from the First image
                xs -- Image points from the Second image

            Return Values:
                A -- Estimated Homography Matrix
        """
        # Construct the DLT Matrix
        A = construct_H_matrix(x, xs)

        # Perform SVD on the Matrix
        U, s, Vh = np.linalg.svd(A.T @ A)

        # Extract the 9th row and Normalize it
        A = Vh[-1, :] / Vh[-1, -1]

        # Reshape the row and get the projection matrix
        A = A.reshape(3, 3)

        return A

In [7]: def find_matching_points(image1, image2, num_points=500):
        """
            Find matching points between two images using SIFT

            Keyword Arguments:
                image1 -- Image 1
                image2 -- Image 2
                num_points -- Maximum number of features
```

```

        Return Values:
        x -- Image points from the First image
        xs -- Image points from the Second image
    """
    # Initiate sift detector
    orb = cv2.xfeatures2d.SIFT_create(nfeatures=num_points)
    # orb = cv2.ORB_create()

    # Create BF Matcher
    # bf = cv2.BFMatcher(cv2.NORM_HAMMING, crossCheck=True)
    bf = cv2.BFMatcher(cv2.NORM_L1, crossCheck=True)

    # Find keypoints and descriptors
    kp1, desc1 = orb.detectAndCompute(image1, None)
    kp2, desc2 = orb.detectAndCompute(image2, None)

    # Match descriptors
    matches = bf.match(desc1, desc2)

    # Sort the matches in the order of their distance
    matches = sorted(matches, key = lambda x:x.distance)[:80]

    # Draw first 10 matches.
    img3 = cv2.drawMatches(image1, kp1, image2, kp2, matches, None, flags=2)

    plt.imshow(img3),plt.show()

    if len(matches) >= 4:
        x = np.array([ kp1[m.queryIdx].pt for m in matches ]).reshape(-1,1,2)
        xs = np.array([ kp2[m.trainIdx].pt for m in matches ]).reshape(-1,1,2)
    return x, xs

In [8]: def warp_image(image_1, image_2, H):
    """
        Warp image_2 to image_1 frame

        Keyword Arguments:
            image_1 -- Image to warp to
            image_2 -- Image to be warped
            H -- 3*3 homography matrix

        Return Values:
            im_out -- Warped image
    """
    im_out = cv2.warpPerspective(image_2, H, (image_1.shape[1] + image_2.shape[1], image_1.shape[0] + image_2.shape[0]), flags=cv2.INTER_LINEAR)
    return im_out

In [9]: def crop_image(image, tolerance=0):

```

```

'''
    Crop out black spaces from images by contouring

    Keyword Arguments:
        image -- Image to be cropped

    Return Values:
        crop -- Cropped image
'''
mask = image > tolerance
if image.ndim==3:
    mask = mask.any(2)
mask0,mask1 = mask.any(0),mask.any(1)
return image[np.ix_(mask1,mask0)]

```

0.3 Image Stitching / Mosaicing

- In order to perform image stitching, we compute the homography between two images using RANSAC.
- After having computed the homography, we transform the second image in the first images reference frame.
- Then, we superimpose both the transformed and the first image in order to create the resultant stitched image.
- In order to stitch multiple images together, we take the resultant image we obtained from the above step and repeat the algorithm
- It is beneficial to perform the process with the resultant image as we get more keypoints which matching a new image and thus relieving us to provide the images in a certain order.

```

In [10]: def stitch_images(images):
'''
    Given a set of images, perform image mosaicing and stitch the images.

    Keyword Arguments:
        images -- Set of images of be stitched

    Return Values:
        A -- Stitched Image
'''
result = images[0]
for i in range(len(images)-1):
    x, xs = find_matching_points(result, images[i+1])
    H = RANSAC(xs, x, len(xs))
    dst1 = warp_image(result, images[i+1], H)

    resultant = np.zeros(dst1.shape).astype('uint8')
    resultant[0:result.shape[0], 0:result.shape[1], :] = result[:, :, :]
    mask = resultant<=0
    resultant[mask] = dst1[mask]

```

```

        result = crop_image(resultant)

        plt.imshow(result)
        plt.show()
    return result

```

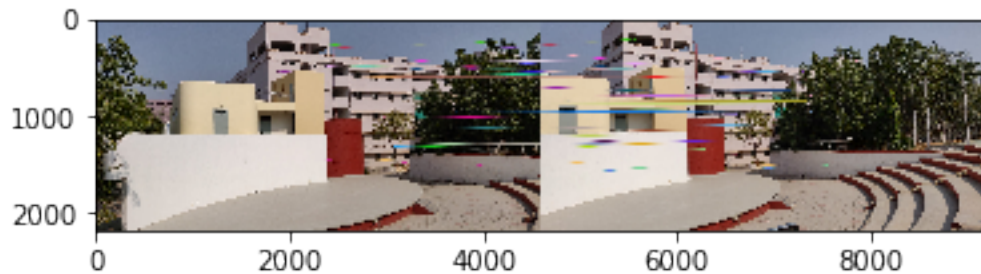
0.3.1 Question 1 : Find matches between two partially overlapped images using any feature detector and descriptor

```

In [11]: image_files = sorted(glob('./image_mosaicing/img1/*'))
        images = []
        for im in image_files:
            images.append(cv2.cvtColor(cv2.imread(im), cv2.COLOR_BGR2RGB))

        x, xs = find_matching_points(images[1], images[2])

```



0.3.2 Question 2 : Estimate homography matrix between two images robustly using RANSAC

```

In [12]: H = RANSAC(xs, x, len(xs))
        print("Estimated Homography Matrix:\n",H)

```

Estimated Homography Matrix:

```

[[ 7.19710665e-01 -1.27012982e-02  1.07937023e+03]
 [-7.09968288e-02  8.79695938e-01  1.45898122e+02]
 [-6.17983387e-05 -5.15009841e-06  1.00000000e+00]]

```

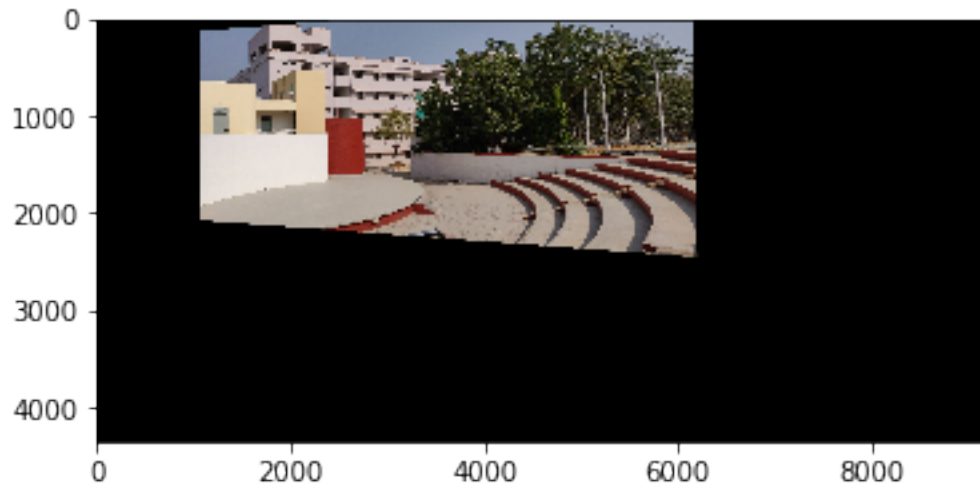
0.3.3 Question 3 : Transform one of the images to the others reference frame homography matrix

```

In [13]: dst1 = warp_image(images[1], images[2], H)
        print("Warped Image")
        plt.imshow(dst1)
        plt.show()

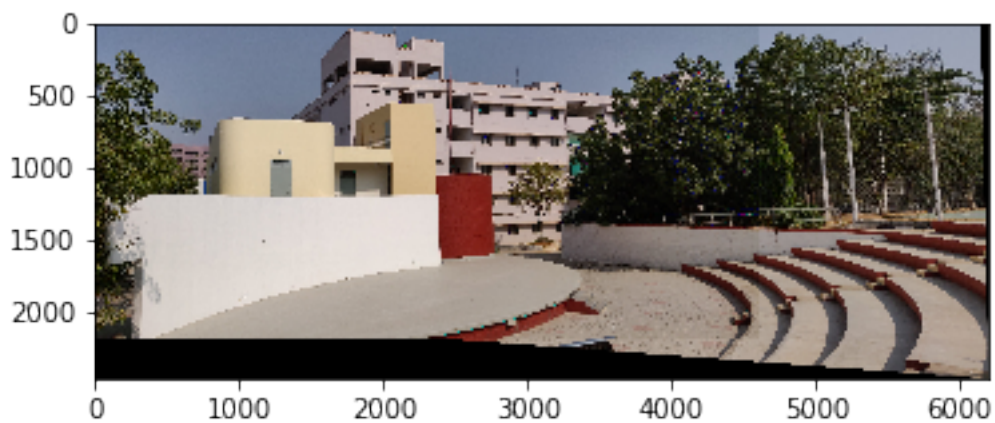
```

Warped Image



0.3.4 Question 4 : Stitch the two images together

```
In [14]: resultant = np.zeros(dst1.shape).astype('uint8')
resultant[0:images[1].shape[0], 0:images[1].shape[1], :] = images[1][:,:, :]
mask = resultant<=0
resultant[mask] = dst1[mask]
resultant = crop_image(resultant)
plt.imshow(resultant)
plt.show()
```

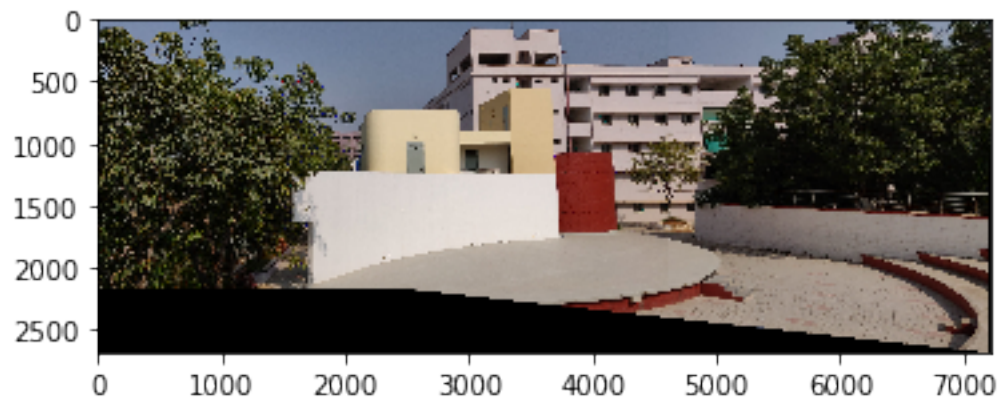


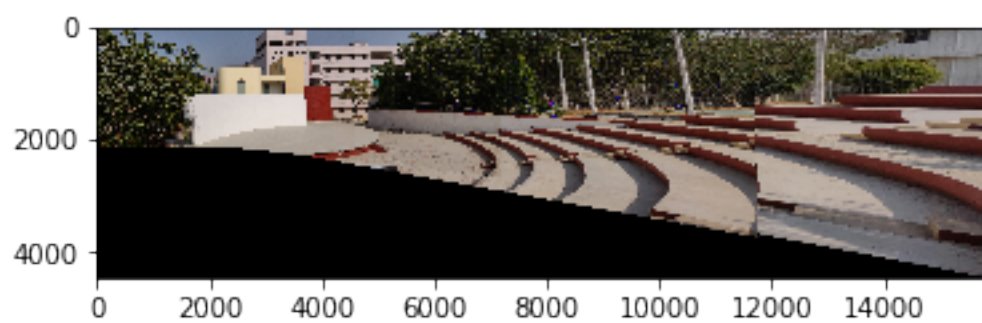
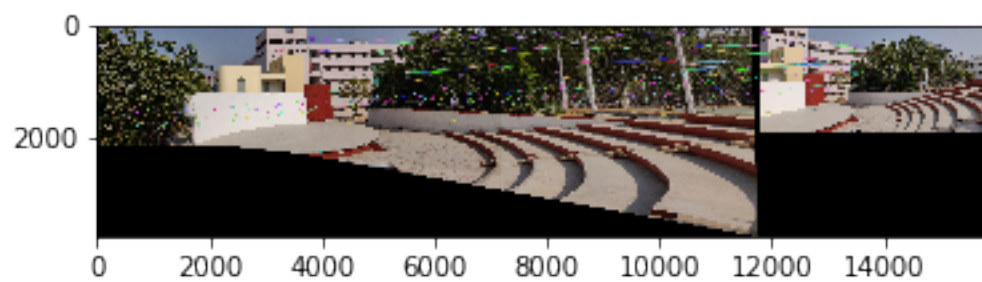
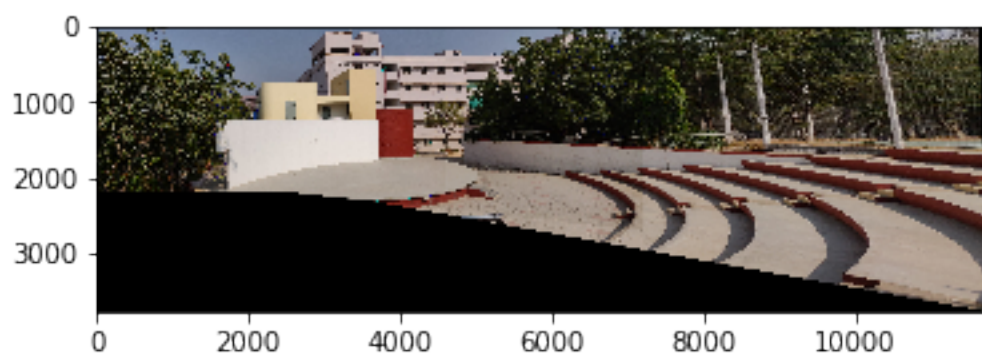
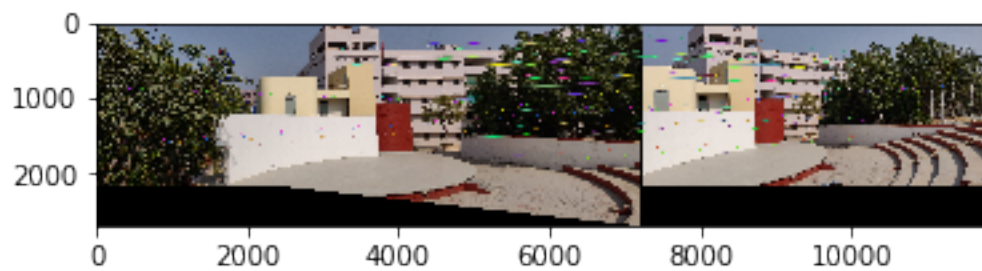
0.3.5 Question 5 : Stitch multiple images together

```
In [17]: image_dirs = sorted(glob('./image_mosaicing/*'))
         stitched_images = []
         for direc in image_dirs:
             print(direc)
             image_files = sorted(glob(direc+'/*'))
             images = []
             for im in image_files:
                 images.append(cv2.cvtColor(cv2.imread(im), cv2.COLOR_BGR2RGB))

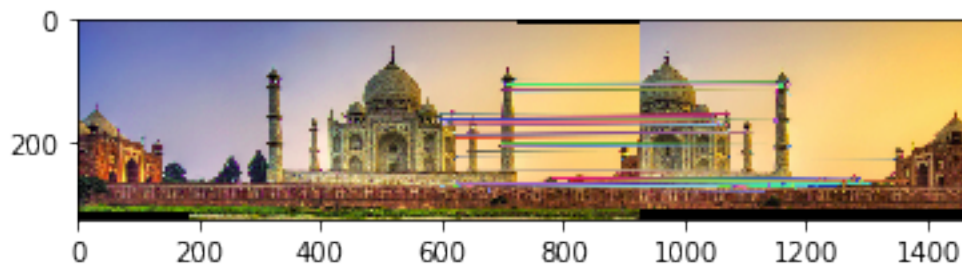
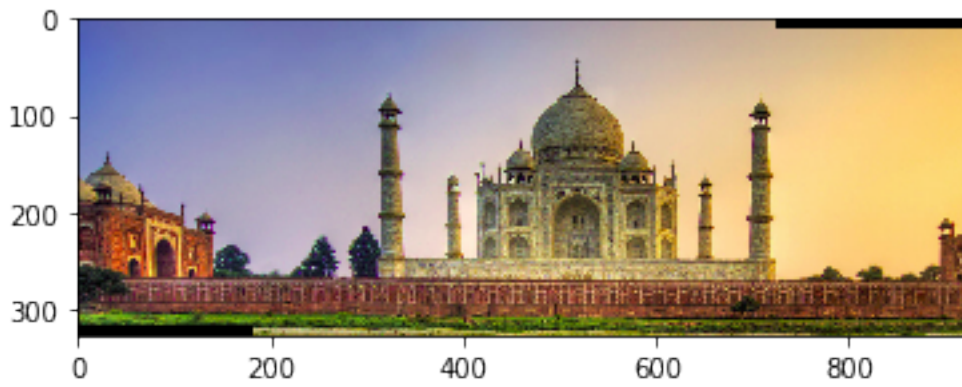
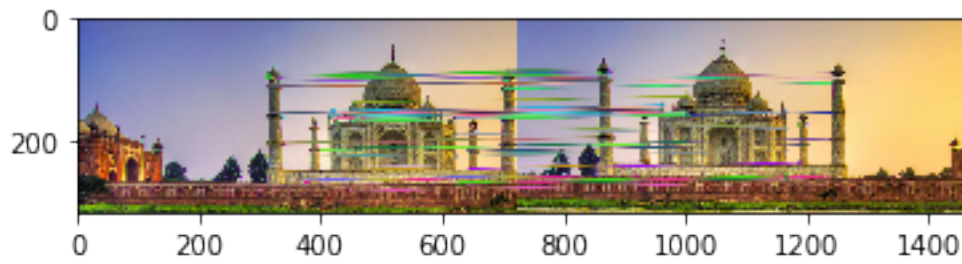
             stitched_images.append(stitch_images(images))

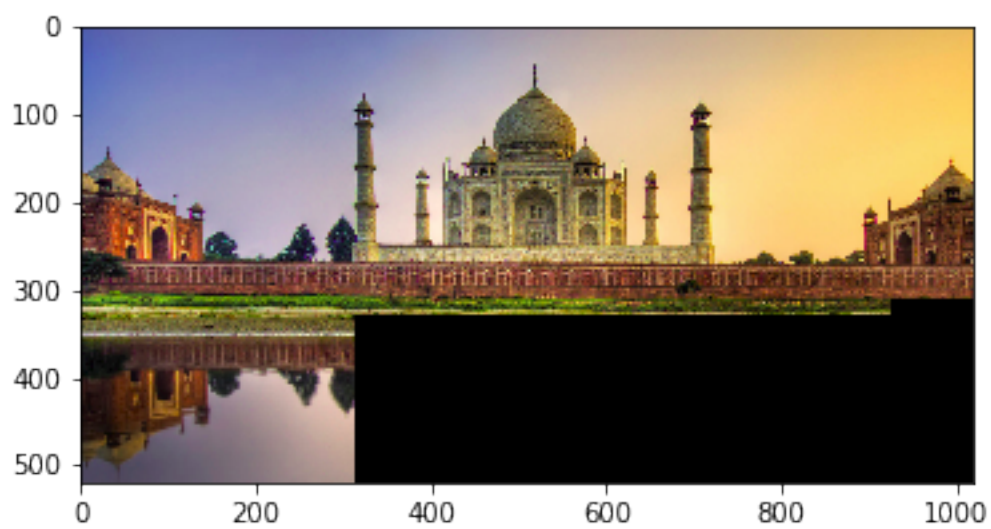
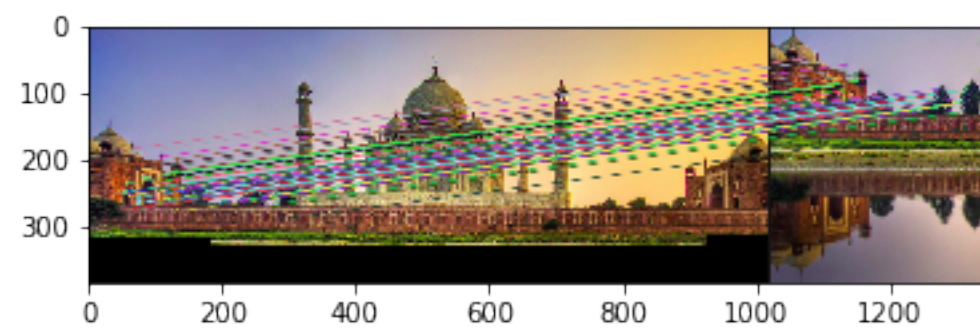
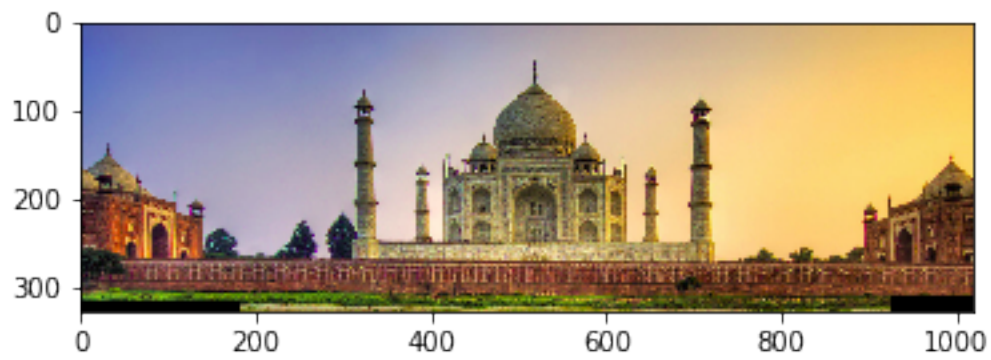
./image_mosaicing/img1
```

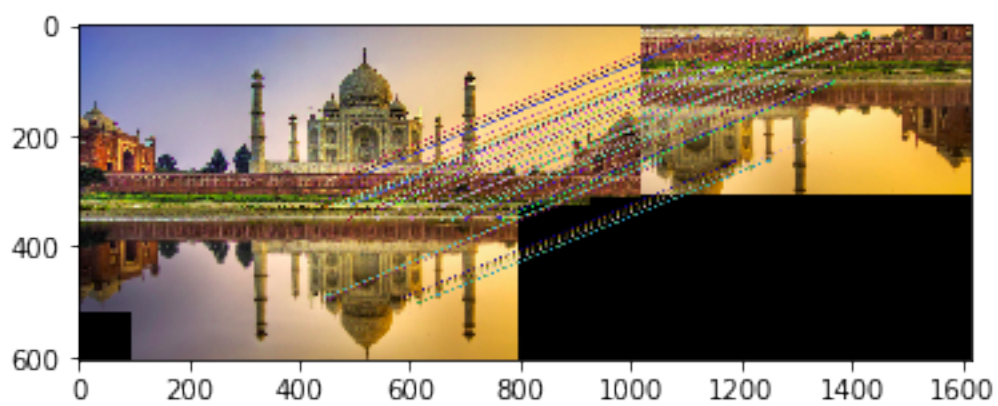
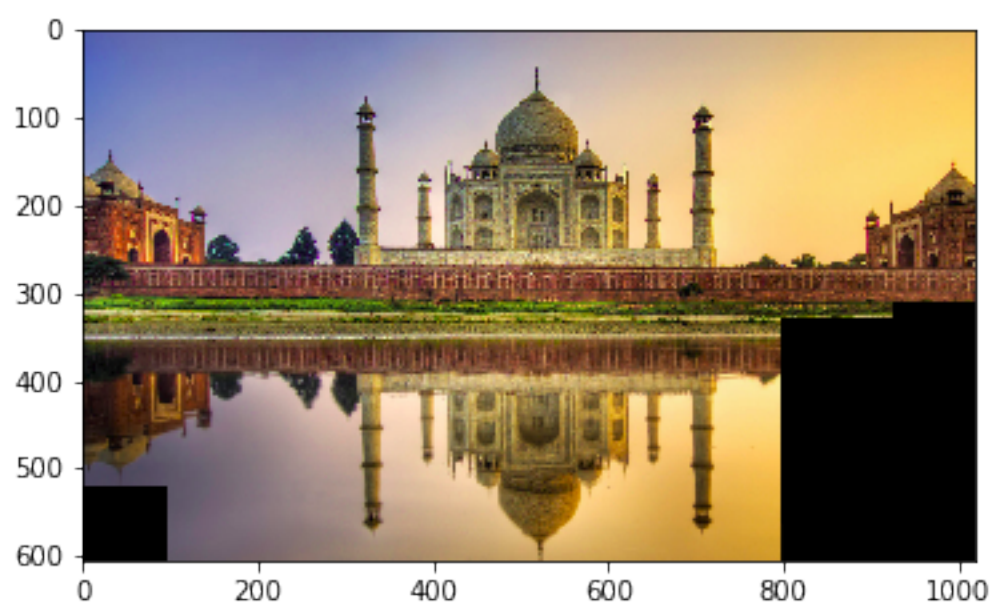
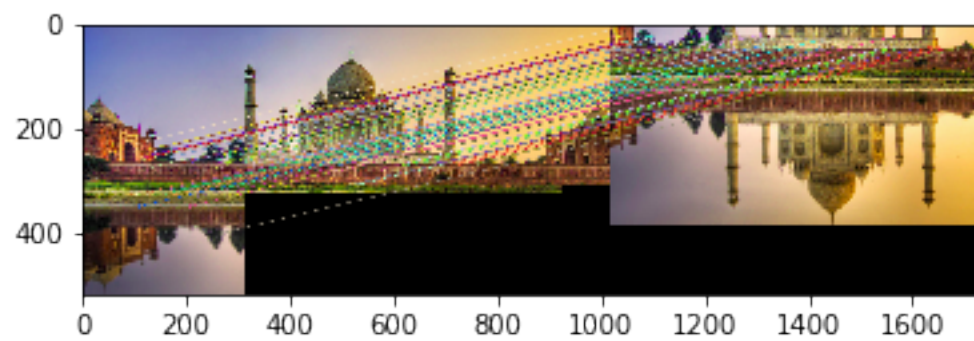


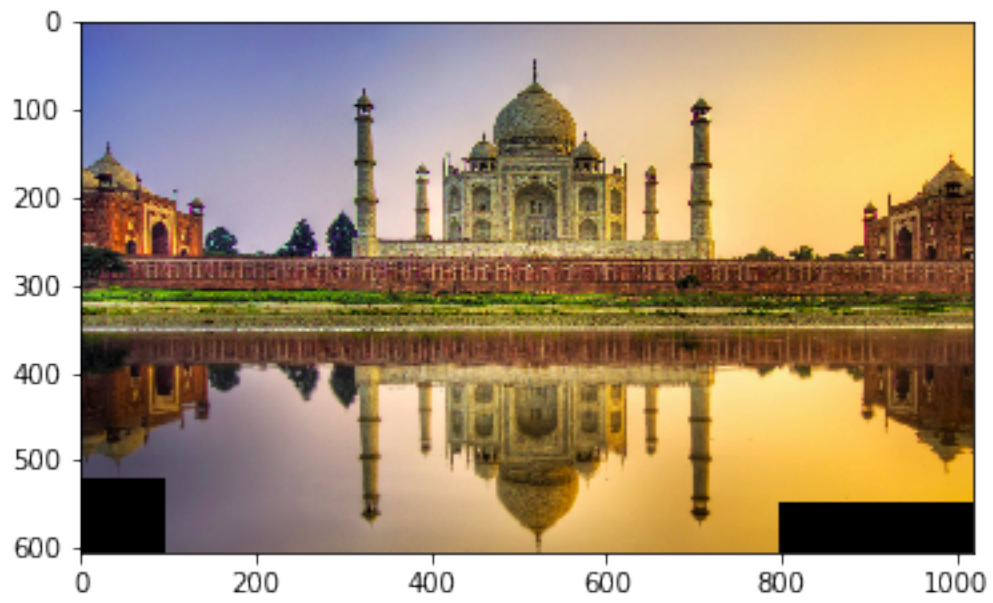


./image_mosaicing/img2

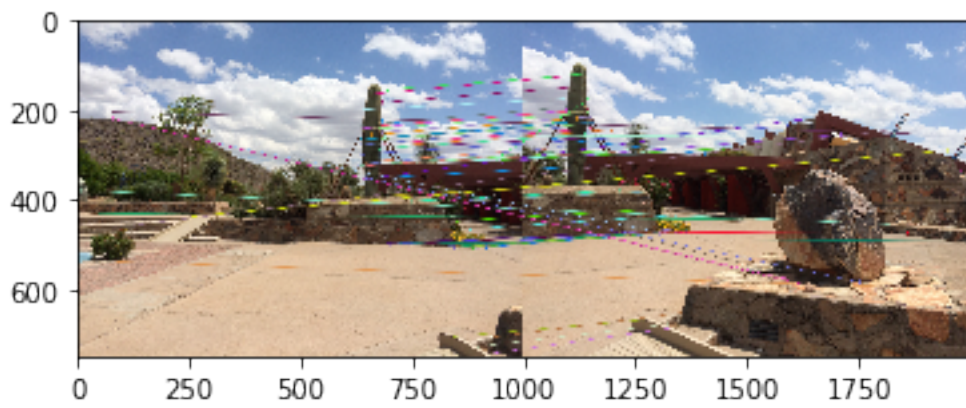


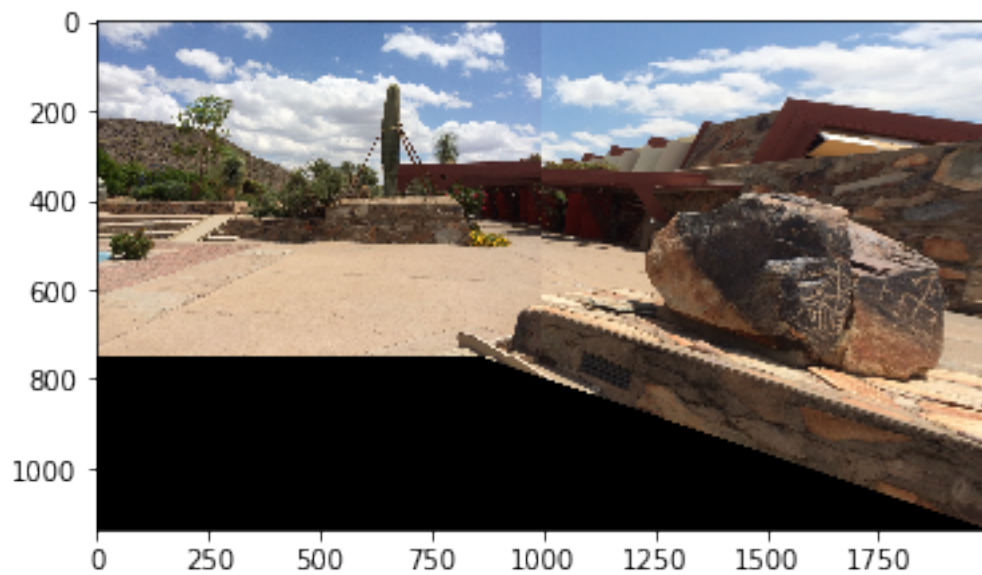




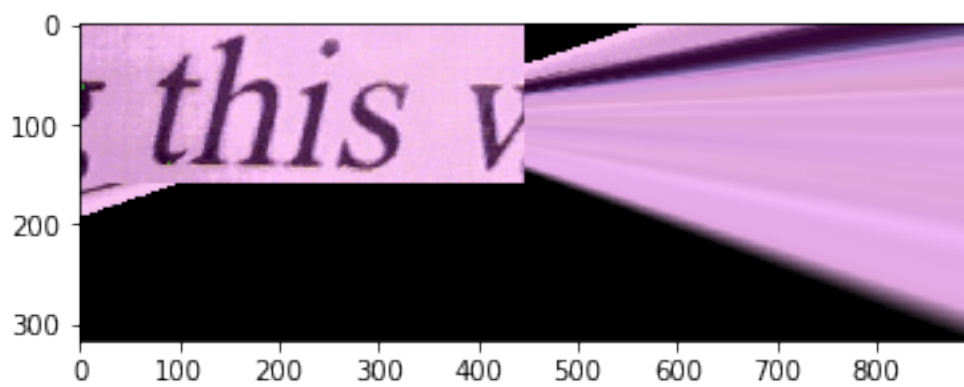
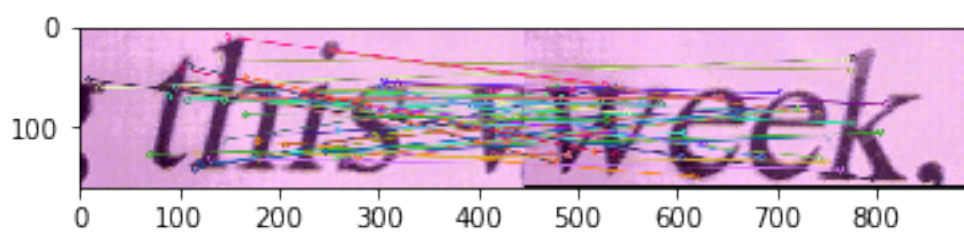


./image_mosaicing/img3

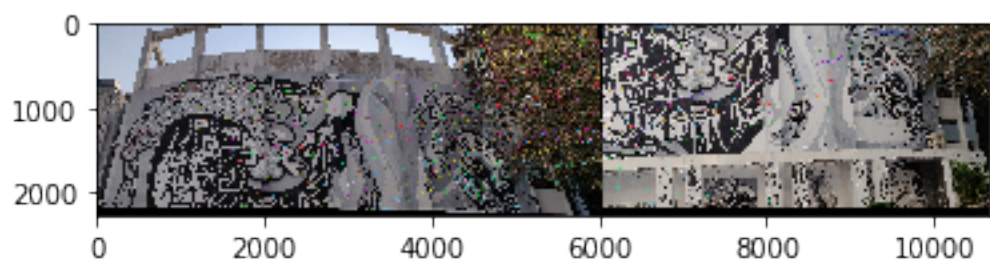
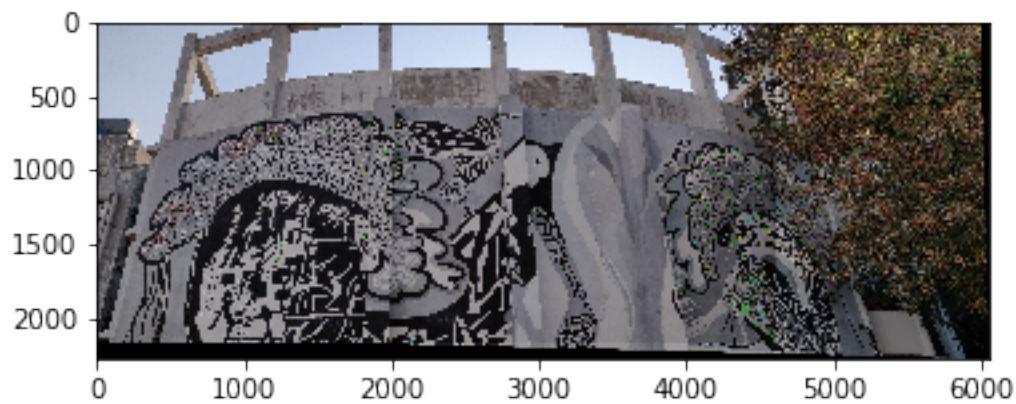
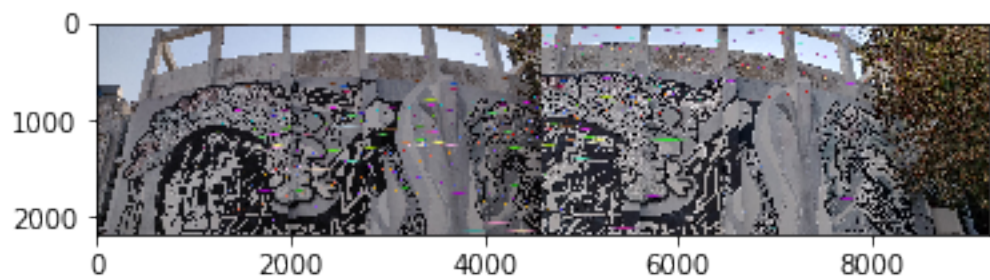


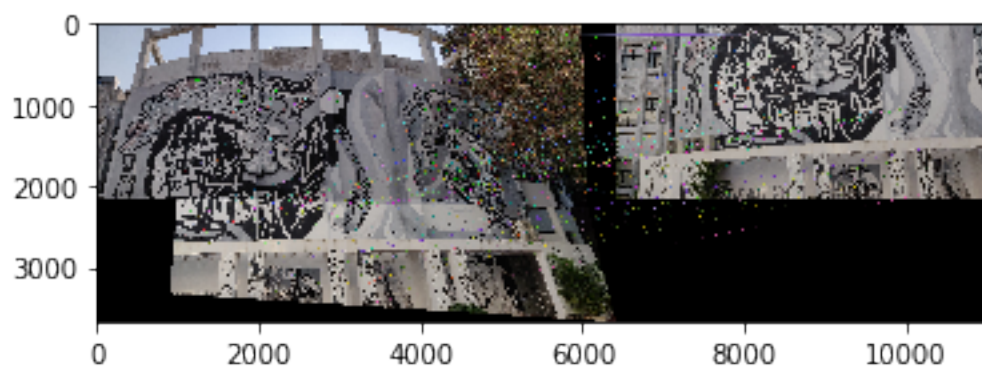
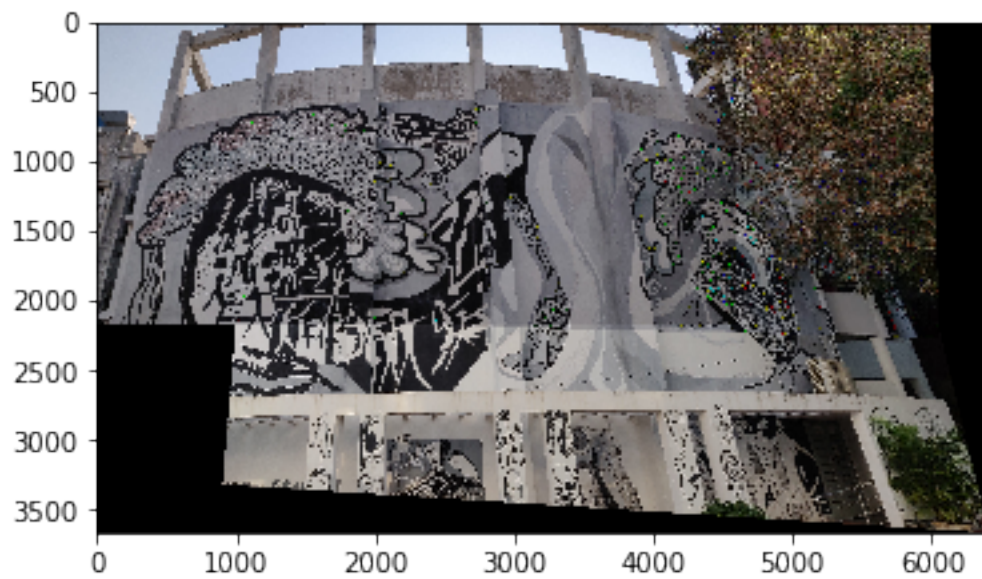


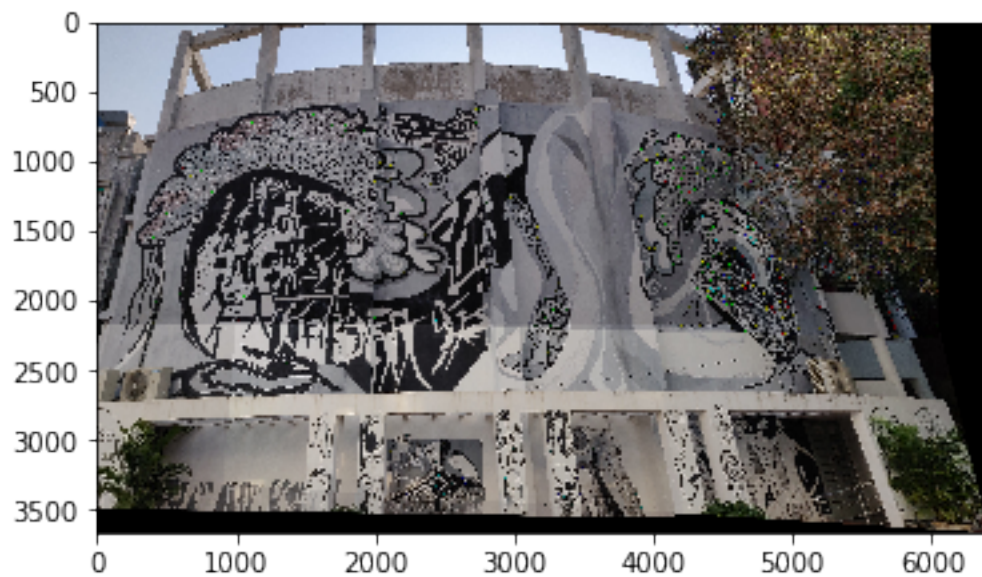
./image_mosaicing/img4



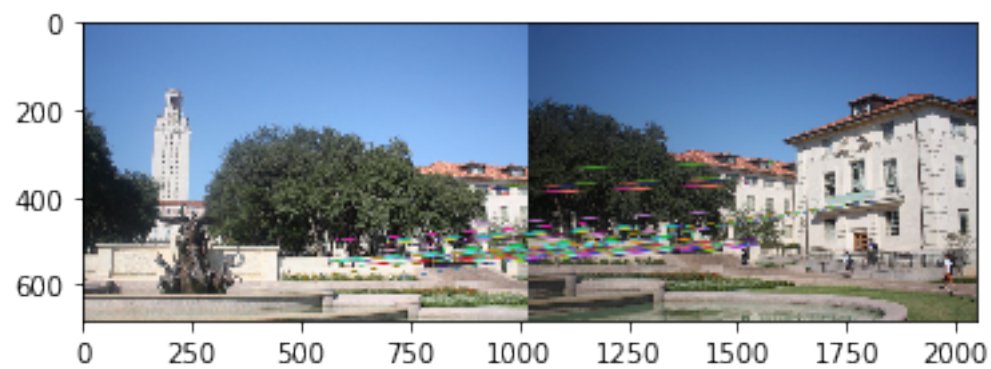
./image_mosaicing/img5

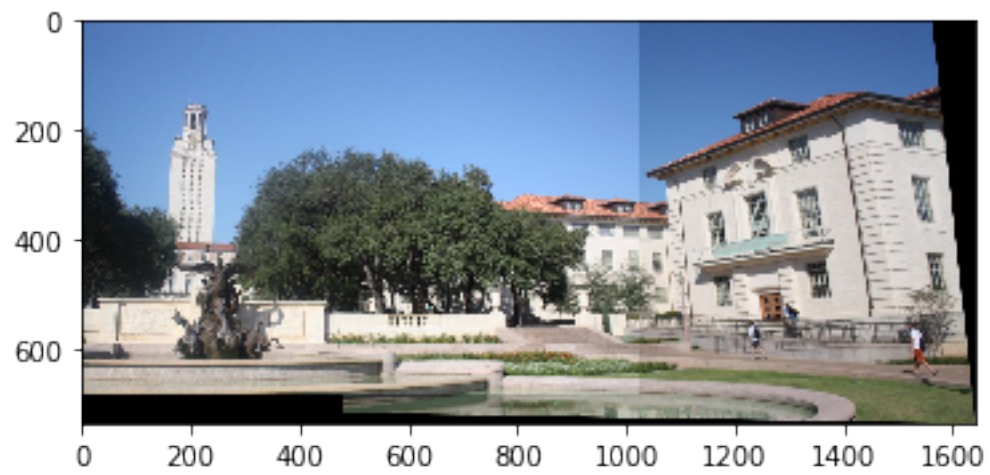




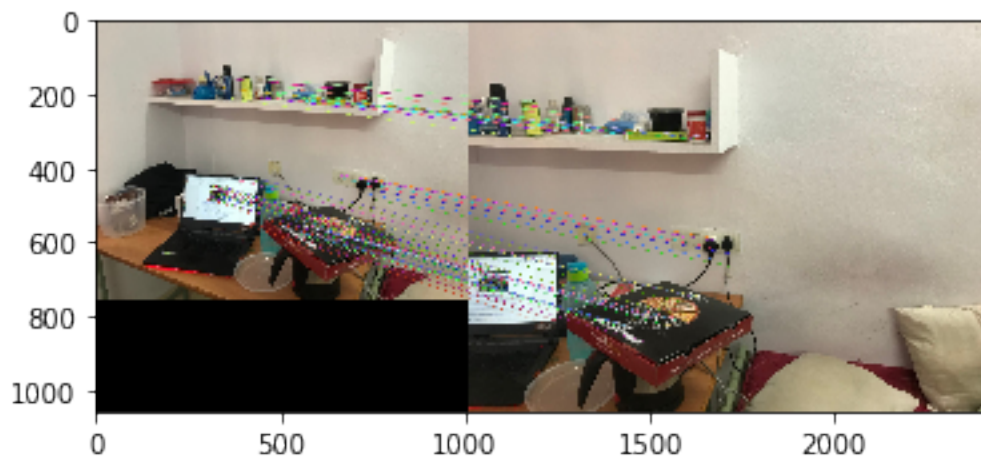


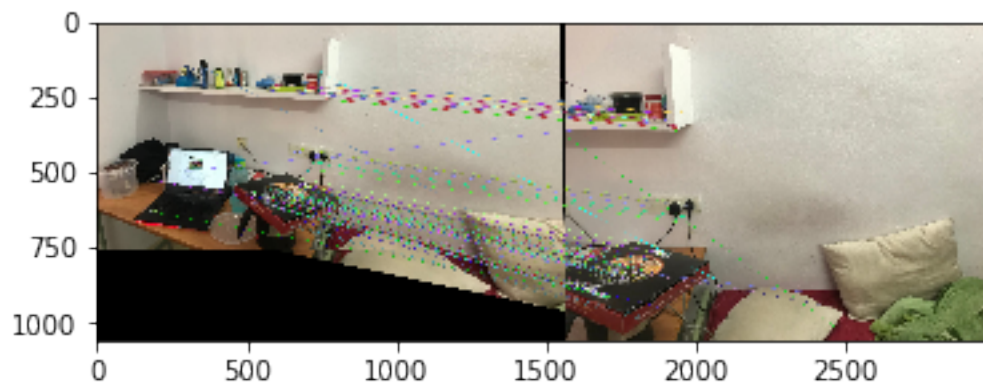
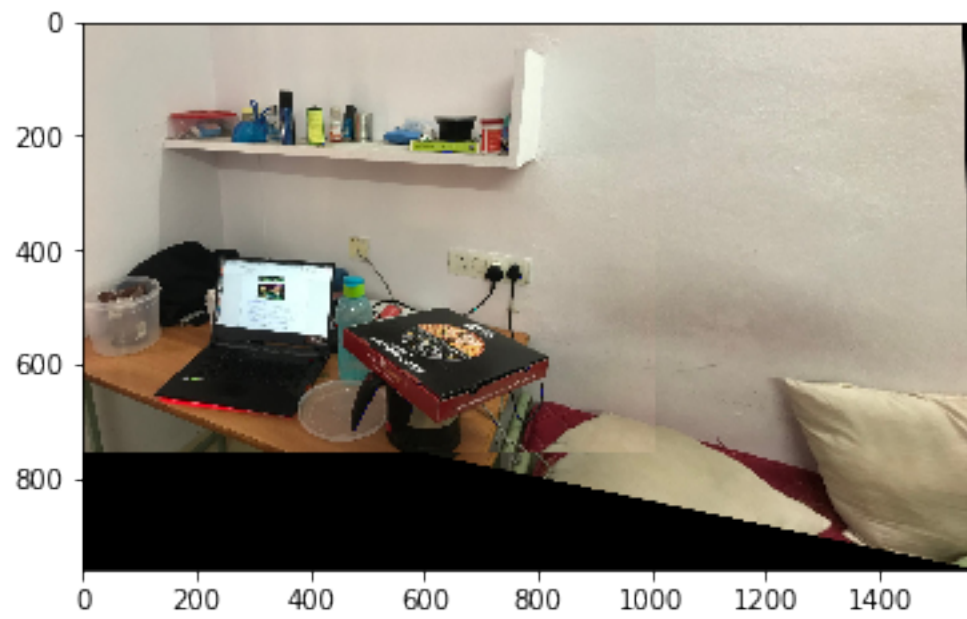
./image_mosaicing/img6

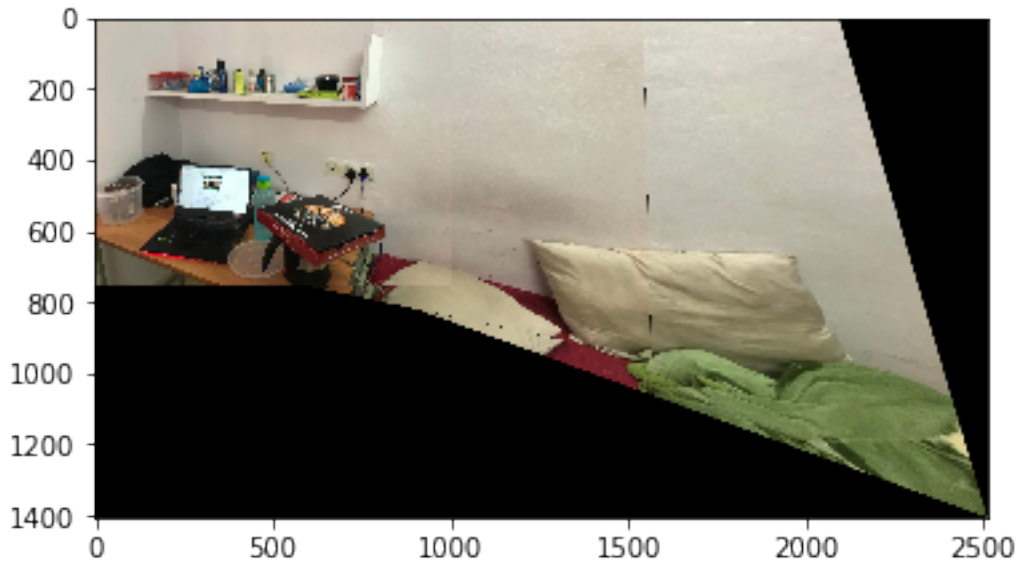




./image_mosaicing/img7







```
In [18]: fig3 = plt.figure(constrained_layout=True,figsize=(20, 20))
gs = fig3.add_gridspec(7,1)
for i in range(len(stitched_images)):
    f3_ax1 = fig3.add_subplot(gs[i])
    f3_ax1.set_title('Left Image')
    plt.imshow(stitched_images[i])
fig3.suptitle('Final Stitched Images', fontsize=16)
plt.show()
```


Final Stitched Images

